

AUTONOMOUS AGENTS · SERIES TWO

The Scholar Series

Lesson 5 — The Human-in-the-Loop Checkpoint.

The most autonomous agent still needs to know when to stop and ask.

★ The capstone — and the responsible end to the whole journey

Where we are. Scholar now loops, plans, reflects, and sustains a long job — it can take a goal and produce a polished report entirely on its own. That power is exactly why this final lesson matters. An agent doing consequential work unsupervised needs a way to *pause and check in* with a human at key moments. This is the responsible counterweight to autonomy, and the right place for both this series and the whole project to end.

What this lesson covers

- 1 **The scene** — an autonomous agent confidently going wrong, unchecked
- 2 **The concept** — pausing for human judgement at the right moments
- 3 **The architecture** — a checkpoint that suspends and resumes the agent
- 4 **Build it in LangGraph** — `interrupt()` and `Command(resume=...)`
- 5 **Build it in CrewAI** — human input on a task
- 6 **Compare, and when to use a checkpoint at all**
- 7 **The whole journey** — two series, the full map of agents

PART 1

The scene: confidently, autonomously wrong

Everything we've built makes Scholar more capable and more independent. But independence has a shadow side. Watch what happens when a fully autonomous Scholar takes a subtly wrong turn early and no one is there to catch it.

Goal: Write an investor briefing on solid-state batteries.

[plan] decides the briefing should focus on consumer electronics applications

[executes all 12 sections around that focus, autonomously, in minutes]

Scholar: Here is your complete 12-section briefing on solid-state batteries in consumer electronics.

The work is polished, internally consistent, and well-reflected — and aimed at entirely the wrong target. The investor actually cared about *automotive* batteries, where the real money is. Scholar made a reasonable-but-wrong framing decision in its very first planning step, then used every capability we gave it to execute that wrong decision beautifully and at high speed. By the time a human sees the output, twelve sections of misdirected work are done. The more autonomous and capable the agent, the more expensive its unchecked mistakes become, because it commits to them faster and more thoroughly.

THE FAILURE, PRECISELY

A fully autonomous agent commits to its own decisions and executes them at speed, with no opportunity for a human to catch a wrong turn before the cost is paid. The capabilities that make Scholar powerful also let it be confidently, efficiently wrong. What's missing is a way to *pause at the right moment* and let a human weigh in before the agent races ahead.

PART 2

The concept: knowing when to ask

The human-in-the-loop checkpoint is the idea that an autonomous agent should *pause at chosen moments* and wait for a human to review, approve, edit, or redirect before continuing. It is not a retreat from autonomy; it is what makes autonomy *safe to deploy*. The agent still does the vast majority of the work on its own — it simply stops at a few high-stakes junctures to get a human's blessing before committing.

The craft lies entirely in *where* you place the checkpoints, and this is a genuine design judgement. Pause too often and you've destroyed the autonomy that made the agent valuable — you're back to driving every step yourself. Pause too rarely and you've let the agent run so far that catching a mistake is expensive. The art is identifying the few moments where a human's judgement is most valuable and a mistake is most costly: typically right after the agent forms its plan (before it executes twelve sections on a wrong premise), and right before any irreversible or consequential action (sending the report, spending money, deleting data).

Notice the beautiful symmetry with where this series began. In Lesson 1, the human gave a goal and then stepped back entirely. Here, in Lesson 5, we invite the human back in — but only at carefully chosen moments, not at every step. We've travelled from "human drives every turn" (series one) through "agent drives itself completely" (Lessons 1 to 4) to the mature middle ground: "agent drives itself, but pauses for the human at the moments that matter." That middle ground is where most real, responsible autonomous systems actually live.


THE ONE IDEA OF THIS LESSON

A human-in-the-loop checkpoint pauses the autonomous agent at chosen high-stakes moments so a human can approve, edit, or redirect before it proceeds. It doesn't reduce autonomy so much as make it deployable. The skill is placing checkpoints at the few junctures where human judgement is most valuable — usually after planning and before irreversible actions.

PART 3

The architecture: a checkpoint that suspends and resumes

Architecturally, a checkpoint is a special point in the agent's flow where execution *suspends* — the agent saves its state and hands control back to a human — and later *resumes* from exactly that point once the human responds. The agent doesn't restart; it picks up precisely where it paused, now carrying the human's decision.



the agent pauses only here — it plans and writes on its own, but checks in before committing

The checkpoint sits after planning. The agent suspends, saving its state; a human approves, edits, or redirects; then the agent resumes from exactly that point with the human's input and runs the rest autonomously.

The two defining features are suspension and resumption. Suspension means the agent genuinely stops and persists its state, rather than blocking or busy-waiting — it can wait minutes or days for the human. Resumption means it continues from the saved checkpoint with the human's response folded in, not from the beginning. This pause-and-continue ability is what lets a human supervise an autonomous agent without babysitting it, which is the entire point.

PART 4

Build it in LangGraph: interrupt and resume

LangGraph has a purpose-built, elegant mechanism for this, and it's the modern recommended approach: the `interrupt()` function pauses the graph inside a node, and `Command(resume=...)` continues it. This relies on the same persistence layer (a checkpointer) that has quietly underpinned LangGraph all along — pausing and resuming is just that persistence put to deliberate use.

```

from typing import TypedDict
from langgraph.graph import StateGraph, END
from langgraph.checkpoint.memory import MemorySaver
from langgraph.types import interrupt, Command
from langchain_openai import ChatOpenAI
from langchain_core.messages import SystemMessage, HumanMessage

llm = ChatOpenAI(model="gpt-4o-mini", temperature=0)

class State(TypedDict):
    goal: str
  
```

```
plan: str
report: str
```

Why this exists: note the imports of `interrupt` and `Command` — the two primitives that pause and resume — and `MemorySaver`, the checkpointer that makes pausing possible.

```
def make_plan(state: State) -> State:
    # The agent writes its plan, as in Lesson 2.
    state["plan"] = llm.invoke([SystemMessage(content="Outline the briefing."),
                               HumanMessage(content=state["goal"])]).content
    return state

def human_checkpoint(state: State) -> State:
    # THE CHECKPOINT. interrupt() pauses the graph here and surfaces
    # the plan to the human. Execution suspends until a human resumes.
    decision = interrupt({
        "question": "Approve this plan, or send edits?",
        "plan": state["plan"],
    })
    # When resumed, 'decision' holds whatever the human sent back.
    if decision != "approve":
        state["plan"] = decision # human edited the plan; use theirs
    return state

def write_report(state: State) -> State:
    state["report"] = llm.invoke([SystemMessage(content=f"Write the full briefing "
        f"following this approved plan: {state['plan']}"),
        HumanMessage(content=state["goal"])]).content
    return state
```

Why this exists: `human_checkpoint` is the whole lesson. `interrupt()` halts the graph and shows the human the plan; the value they send back becomes `decision`. The agent only writes the report after the human has weighed in.

```
# A checkpointer is REQUIRED – it persists state across the pause.
graph = StateGraph(State)
graph.add_node("plan", make_plan)
graph.add_node("checkpoint", human_checkpoint)
graph.add_node("write", write_report)
graph.set_entry_point("plan")
graph.add_edge("plan", "checkpoint") # pause AFTER planning
graph.add_edge("checkpoint", "write") # resume INTO writing
graph.add_edge("write", END)
agent = graph.compile(checkpointer=MemorySaver())

# Run until the interrupt; the graph pauses and returns control.
config = {"configurable": {"thread_id": "report-1"}}
agent.invoke({"goal": "investor briefing on solid-state batteries"}, config)

# ...human reviews the plan, then we RESUME with their decision:
agent.invoke(Command(resume="approve"), config) # or resume with edits
```

Why this exists: the `thread_id` lets LangGraph find the paused state to resume. The first `invoke` runs to the checkpoint and stops; the second, with `Command(resume=...)`, continues from exactly that point. That's suspend-and-resume in action.

PART 5

Build it in CrewAI: human input on a task

CrewAI offers a refreshingly simple version for the common case: set `human_input=True` on a task, and CrewAI will pause after that task to let a human review and give feedback before the crew continues. It's the checkpoint idea distilled to a single flag — less flexible about *where* the pause lands, but remarkably easy for the standard "review this output" case.

```
from crewai import Agent, Task, Crew, Process

planner = Agent(role="Planner", goal="Outline the briefing",
                backstory="You map out coverage before writing.")
writer = Agent(role="Writer", goal="Write the briefing from the approved plan",
               backstory="You follow the approved outline faithfully.")

# human_input=True makes CrewAI PAUSE after this task for human review.
plan_task = Task(
    description="Outline the investor briefing on solid-state batteries.",
    expected_output="A section-by-section outline.",
    agent=planner,
    human_input=True,    # ← the checkpoint, in one flag
)
write_task = Task(
    description="Write the briefing following the approved outline.",
    expected_output="The full briefing.",
    agent=writer, context=[plan_task],
)

crew = Crew(agents=[planner, writer], tasks=[plan_task, write_task],
            process=Process.sequential)
result = crew.kickoff()    # pauses for your input after planning
```

Why this exists: `human_input=True` is CrewAI's entire checkpoint mechanism. After the planning task, the crew stops and asks for your review before the writer proceeds — the same suspend-and-resume idea, with almost no code.

PART 6

Compare, and when to checkpoint at all

One last framework comparison, and then a broader point that applies no matter which tool you choose.

LANGGRAPH VS CREWAI — HUMAN-IN-THE-LOOP

	LangGraph	CrewAI
Mechanism	interrupt() anywhere + Command(resume)	human_input=True on a task
Placement	Any point you choose, fully flexible	After a task — simpler, less flexible
Resume	Explicit, with structured human data	Handled by the framework prompt
Effort	More setup (checkpointer, thread_id)	One flag
Best when	You need precise, custom checkpoints	Standard "review this output" pauses

The framework verdict mirrors the whole series: LangGraph gives you precise control over exactly where and how the agent pauses, at the cost of more setup, while CrewAI gives you a one-flag checkpoint for the common case, at the cost of flexibility about placement. But there's a more important judgement that sits above the tool choice, and it's the right note to end the series on: *not every agent needs a human in the loop, and not every step needs a checkpoint*. The decision to add one should come from the stakes. Low-stakes, easily-reversible work can run fully autonomously. High-stakes, costly, or irreversible work earns a checkpoint. Mature engineering is matching the degree of oversight to the degree of risk — neither babysitting a harmless agent nor letting a consequential one run unchecked.

PART 7

The whole journey: two series, the full map

You've reached the end of both series, so let's stand back and see the whole landscape you've now crossed. Two series, ten lessons, and together they cover the two great families of agent design.

The ShopHub series taught the *reactive* agent — one that waits for a request and responds — growing it one capability at a time:

1 Conversation	held a chat by replaying history
2 Memory	recognised returning customers across sessions
3 Routing	classified each message to a specialist (LangGraph)
4 Tools	fetches live facts via search, then MCP
5 A team	a coordinator orchestrating specialist agents

The Scholar series taught the *autonomous* agent — one that pursues a goal on its own — and added the skill of choosing between frameworks:

1	Agentic loop	think-act-observe, driving itself (LangGraph + CrewAI)
2	Self-planning	wrote its own roadmap before working
3	Reflection	critiqued and revised its own output
4	Progress memory	tracked its own state through a long job
5	Human-in-the-loop	paused for human judgement at key moments

WHAT YOU'VE ACTUALLY LEARNED, ACROSS BOTH SERIES

You've seen the two halves of agent design and how they complement each other. A *reactive* agent answers when asked; an *autonomous* agent goes off and does a job. Real-world systems blend the two — and you now have the mental model to recognise the parts of any agent you meet, however complex, plus the judgement to choose a framework and decide how much human oversight a task deserves. That last part — matching the tool and the oversight to the job — is what separates someone who has read about agents from someone who can responsibly build them.

The concepts you arrived with as vocabulary — loops, planning, reflection, memory, tools, orchestration, human-in-the-loop — are now things you've built twice over, watched run, and seen expressed in more than one framework. You can look at any agent system and see its skeleton. And just as importantly, you can ask the mature questions: does this need to be autonomous or reactive, which framework fits, and where does a human belong in the loop? That is the real destination of both series.

End of Scholar Lesson 5, and the end of both series. Scholar grew from a single loop into a responsible, supervisable autonomous agent — and you grew from reading about agents into being able to design, build, and govern them. Now go build something of your own.